

# Memory Error Detect Enhancement for DPDK Fuzzing

XUEQIN LIN, YINAN WANG, ZHAOYAN CHEN

INTEL

# Agenda



- Memory Detect Method
- Overview for Address Sanitizer, method
- Enable Asan in DPDK
- Why DPDK Fuzzing
- DPDK libFuzzer Deployment
- Summary & Next plan

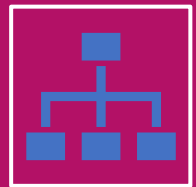
# Background



Hard debug



Critical impact



Insufficient  
management for DPDK  
RTE\_MALLOC\_DEBUG

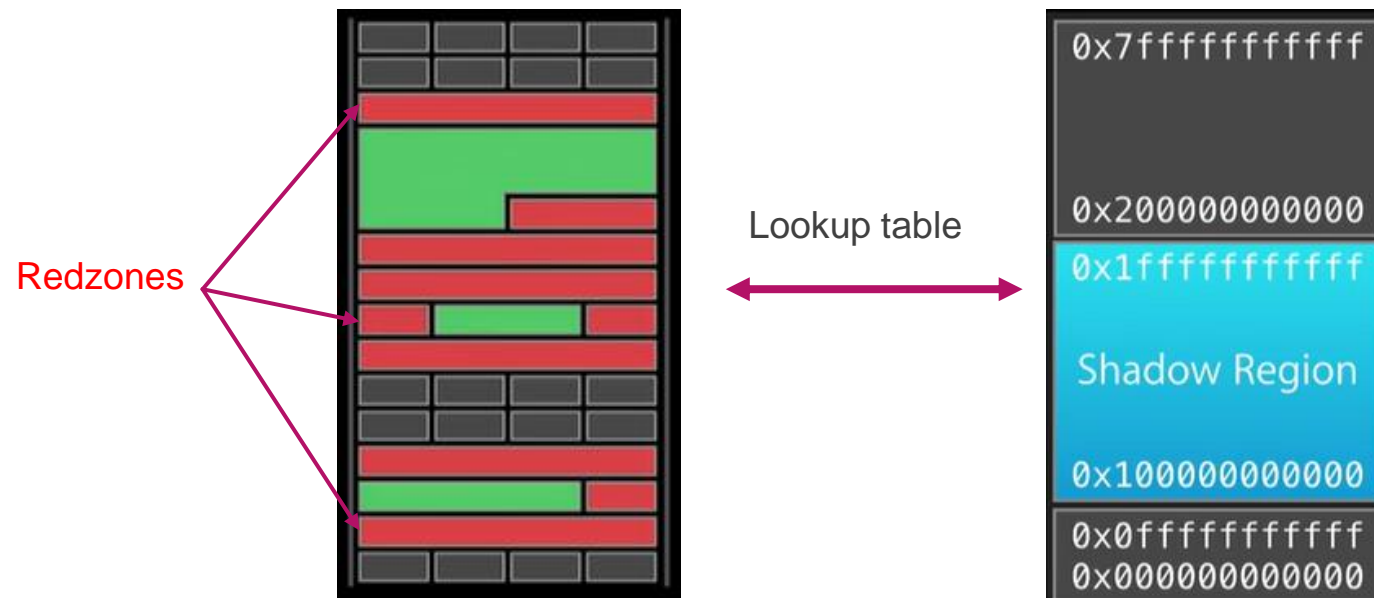


No memory tool could  
use directly on DPDK

# Memory detect method

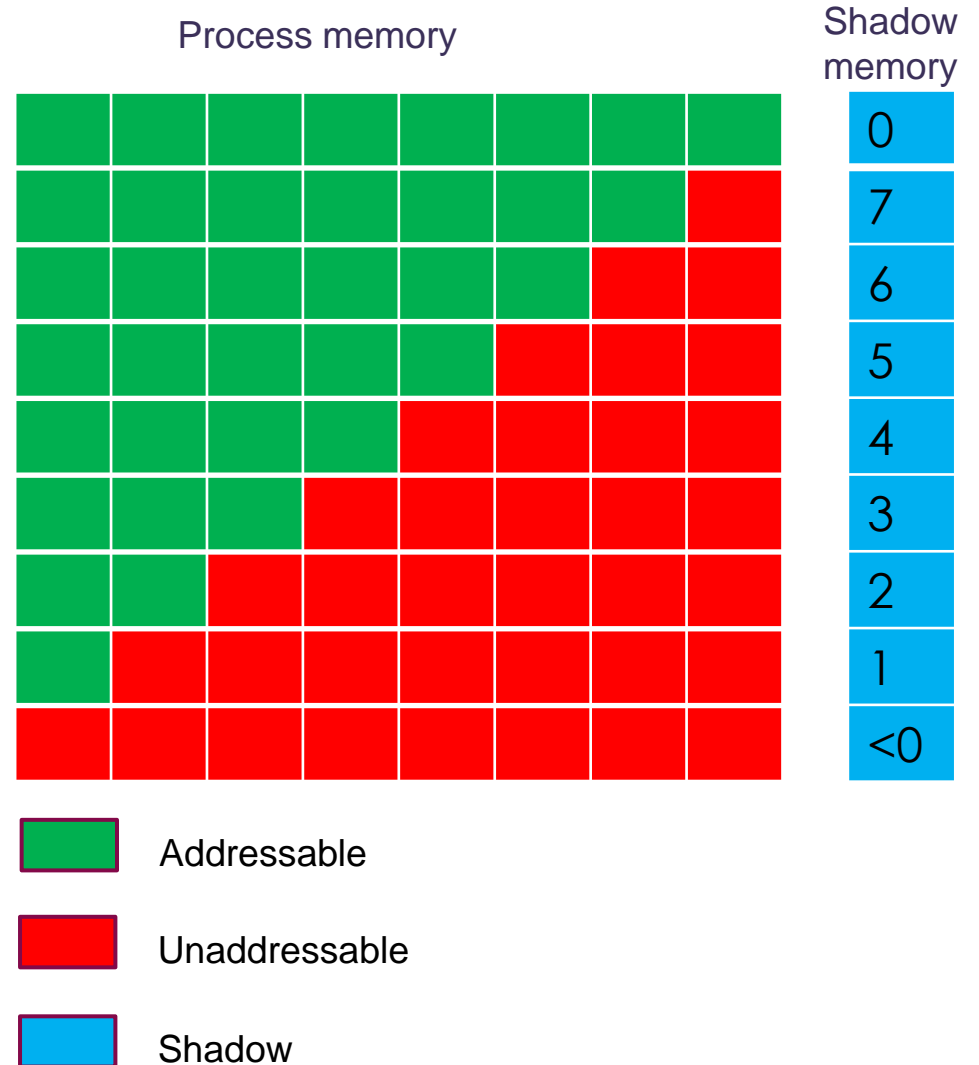


- ❑ Create poisoned **redzones** around memory allocation
- ❑ Instrument all loads/stores to check **shadow state** for each memory access
  - Build **lookup table**
- ❑ **Run-time** detect and **bookkeeping** for error messages

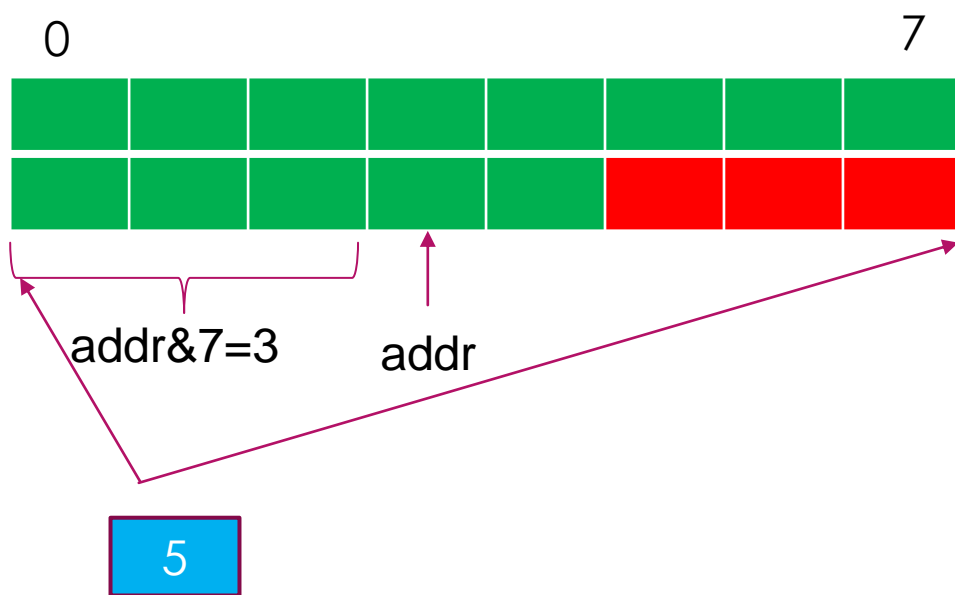


# Shadow byte

- ❑ Every aligned 8-byte word of memory has 9 states
- ❑ All of 8 bytes are addressable, shadow state is 0
- ❑ All of 8 bytes are unaddressable, shadow state is negative(<0), start with “ f ”, use “ fa ” for heap, “ f1 ” for stack...
- ❑ First N bytes are addressable, the rest 8-N bytes are not( $1 \leq N \leq 7$ )
- ❑  $\text{ShadowAddr} = (\text{Addr} \gg 3) + \text{Offset}$



# Instrumentation



## ❑ 8 bytes access

```
long *addr = (long *)0xffff800012345678;
```

```
char *shadow = (char *)(((unsigned long)addr >> 3) +  
SHADOW_OFFSET);  
if (*shadow)  
    report_bug();
```

```
*addr = 0;
```

## ❑ 1, 2, 4 bytes access

```
//replace above if() as below, N=1,2,4
```

```
if (*shadow && *shadow < ((unsigned long)addr & 7) + N)
```

# Address Sanitizer overview

## Support

- LLVM 3.1 (GCC, Clang)
- Support Linux i386/x86\_64, MacOS, Android, Windows...

## Detect

- Buffer overflows(heap, stack, global variable)
- Stack-use-after-return, heap-use-after-free
- more

# Run-time library



## Initialize shadow memory at startup

- Shadow mapping on Linux/x86\_64 with `SHADOW_OFFSET = 0x00007fff8000`

## Provide full malloc/free replacement

- Insert poisoned redzones around allocated memory
- Quarantine for free-ed memory

## Collect stack traces for every malloc/free

## Bookkeeping for error messages



# Enable Asan in DPDK



Why DPDK can't use Asan directly?

DPDK use  
different memory  
management and  
API from glibc

Run-time library  
libasan.so can't  
hook DPDK  
memory API

Can't set redzone  
and add to  
shadow mapping

# Quick enable Asan in DPDK



## ❑ Provide malloc/free replacement

- ❑ Enable RTE\_MALLOC\_DEBUG, add trailer cookie(redzone) in heap\_alloc() to help identify buffer overflows
- ❑ Calculate redzone and shadow address, map redzone and shadow memory address

```
#define MALLOC_ELEM_TRAILER_ADDR(elem) (((void *)RTE_PTR_ADD(elem, elem->size - MALLOC_ELEM_TRAILER_LEN)))
```

```
#define MEM_TO_SHADOW(mem) (((mem) >> 3) + 0x00007fff8000)
```

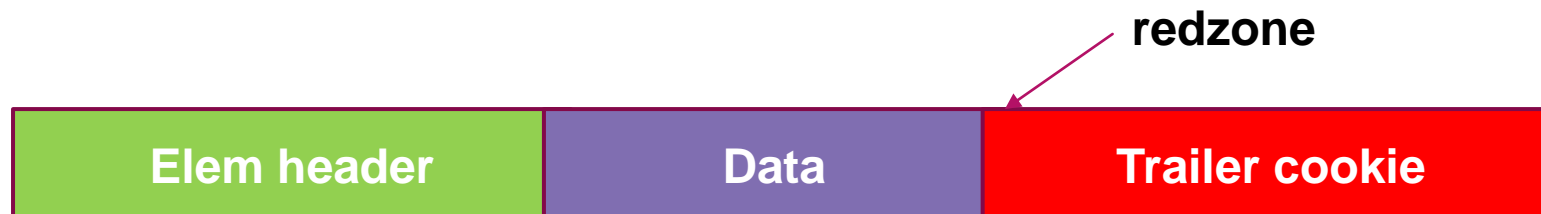
```
...
```

```
int64_t ptr = (uint64_t)MALLOC_ELEM_TRAILER_ADDR(elem);
```

```
char *shadow = (char *)MEM_TO_SHADOW(ptr);
```

- ❑ Set redzone as unaddressable, "fa" as heap negative state.

```
set_redzone(shadow, 0xfa);
```



# Quick enable Asan in DPDK



- ❑ Add redzone in malloc\_heap\_free() for free

```
uint64_t ptr = (uint64_t)MALLOC_ELEM_TRAILER_ADDR(elem);  
char *shadow = (char *)MEM_TO_SHADOW(ptr);
```

- ❑ **Support run-time library libdpdk\_asan.so build in meson.build**

```
add_project_link_arguments('-ldpdk_asan', language: 'c')  
dpdk_extra_ldflags += '-ldpdk_asan'
```

- ❑ **Compile code**

- ❑ Compile with -Db\_sanitize=address flag to enable ASan tool
- ❑ Compile with -Dbldtype=debug to enable gdb for more debug information

# Hello-world sample: heap-buffer-overflow



```
char *p = rte_zmalloc(NULL, 64 , 64);  
p[65] = 'a';
```

```
=====
hello from core 3
==104064==ERROR: AddressSanitizer: heap-buffer-overflow on address 0x7f6b3f9589c1 at pc 0x562e1c9d6578 bp 0x7ffc611f1b40 sp 0x7ffc611f1b30
WRITE of size 1 at 0x7f6b3f9589c1 thread T0
    #0 0x562e1c9d6577 in main ../examples/helloworld/main.c:44
    #1 0x7f731d8abb96 in __libc_start_main (/lib/x86_64-linux-gnu/libc.so.6+0x21b96)
    #2 0x562e1c9d6369 in _start (/root/memory_debug/dpdk_asan311/x86_64-native-linuxapp-gcc/examples/dpdk-helloworld+0x6fa369)

SUMMARY: AddressSanitizer: heap-buffer-overflow ../examples/helloworld/main.c:44 in main
Shadow bytes around the buggy address:
  0x0fede7f230e0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x0fede7f230f0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x0fede7f23100: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x0fede7f23110: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x0fede7f23120: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
=>0x0fede7f23130: 00 00 00 00 00 00 00 00[fa]00 00 00 00 00 00 00 00
  0x0fede7f23140: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x0fede7f23150: 00 00 00 00 00 00 00 00 fa 00 00 00 00 00 00 00
  0x0fede7f23160: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x0fede7f23170: 00 00 00 00 00 00 00 00 fa 00 00 00 00 00 00 00
  0x0fede7f23180: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
Shadow byte legend (one shadow byte represents 8 application bytes):
Addressable:          00
Partially addressable: 01 02 03 04 05 06 07
Heap left redzone:    fa
Freed heap region:    fd
Stack left redzone:    f1
Stack mid redzone:    f2
```

# Agenda

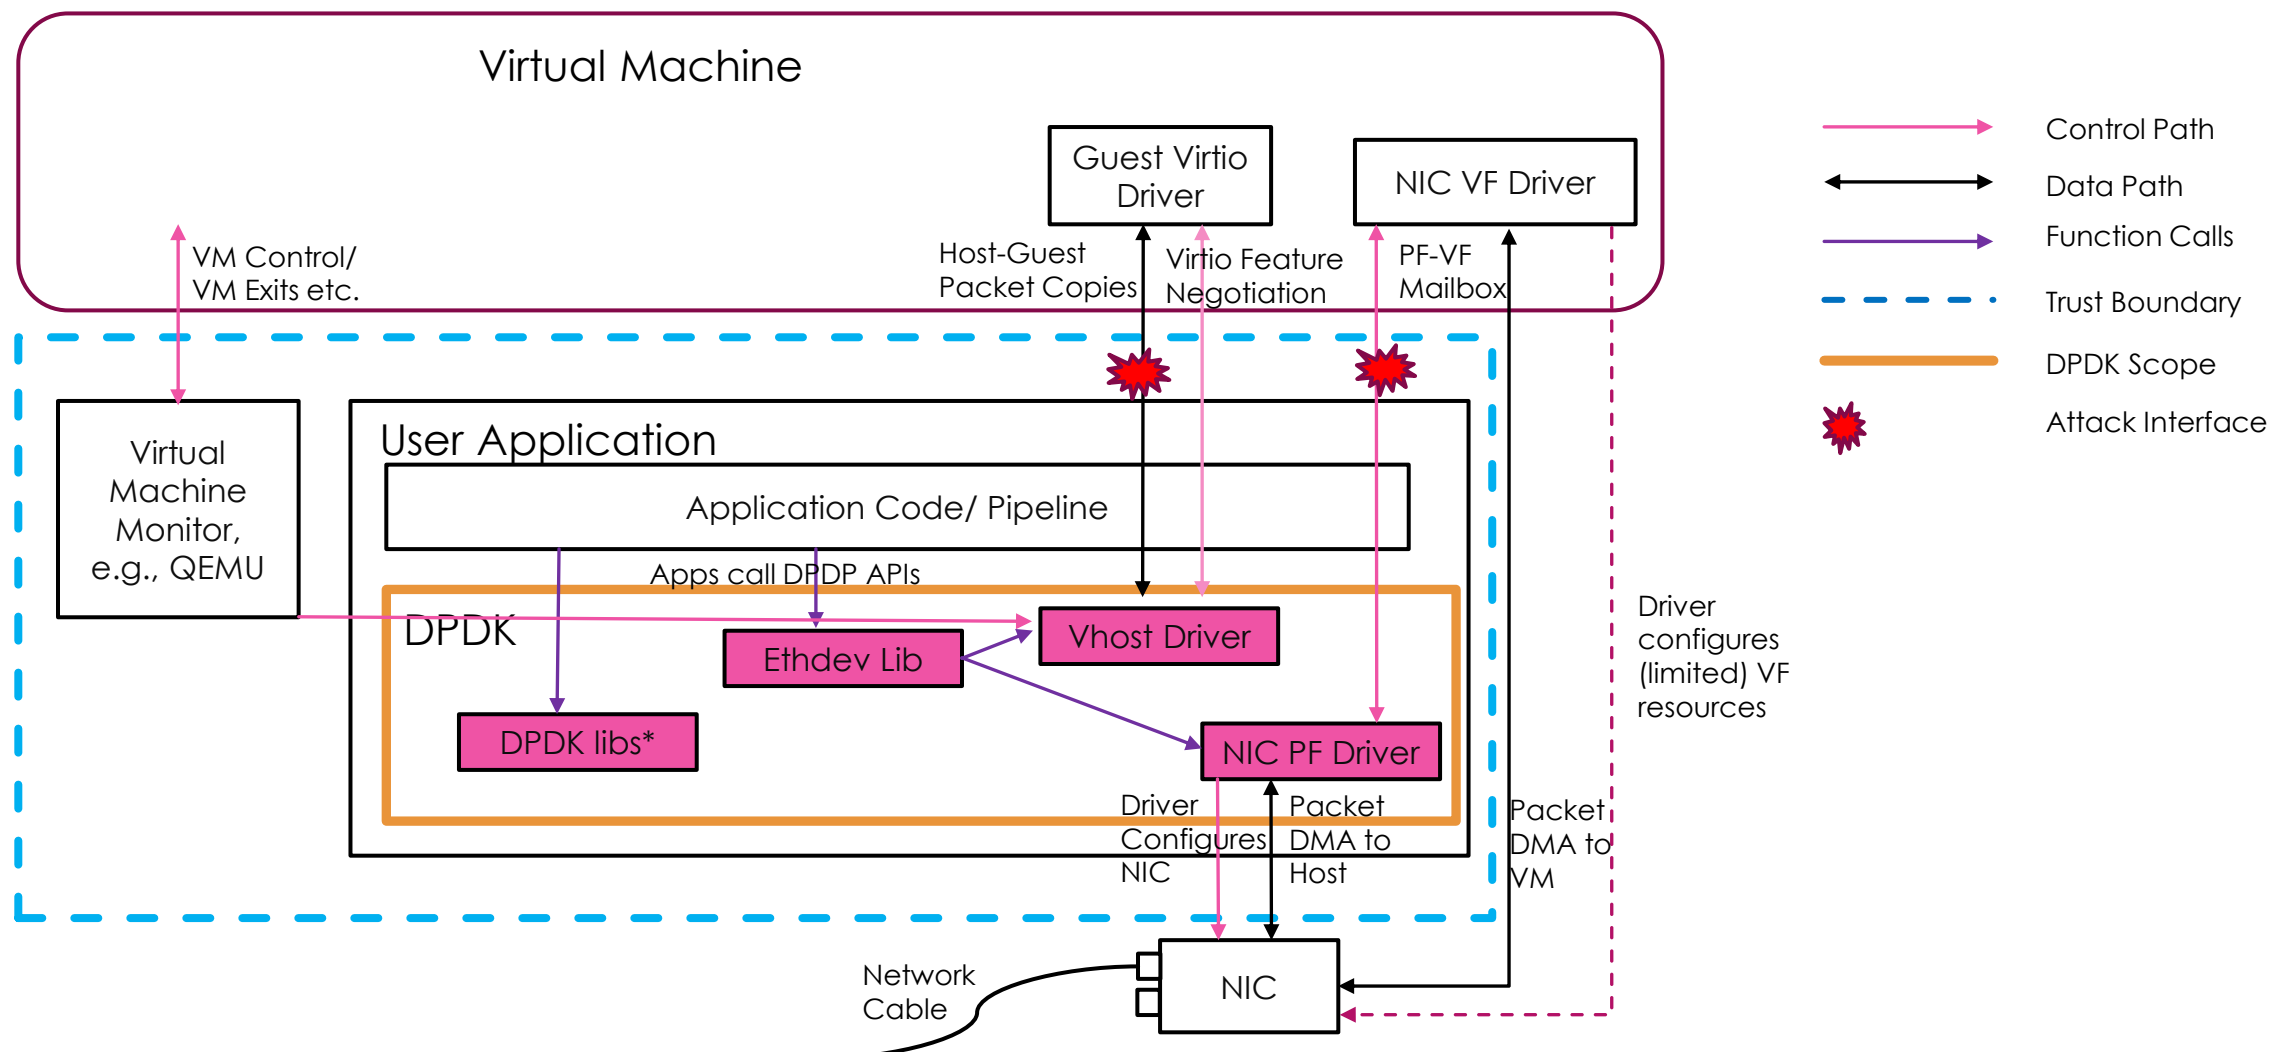


- ❑ Background
- ❑ Memory Detect Method
- ❑ Address Sanitizer Overview
- ❑ Enable Asan in DPDK
- ❑ **Why DPDK Fuzzing**
- ❑ **DPDK libFuzzer Deployment**
- ❑ **Summary & Next plan**

# Why DPDK Fuzzing



## DPDK Threat Modeling



# Why DPDK Fuzzing



- ❑ Memory related uncovered issues are typically vulnerabilities that can be exploited, e.g. **CVE-2020-14374\***
- ❑ Identify **attack interface** memory related issues is vital important
- ❑ Fuzzing can help to finds bugs that are overlooked by regular scripted testing
- ❑ Fuzzing can utilize **detect tools** provided by compiler , e.g., Address Sanitizer

\* Note: DPDK vhost多个安全漏洞通告-启明星辰 ([venustech.com.cn](https://www.venustech.com.cn))

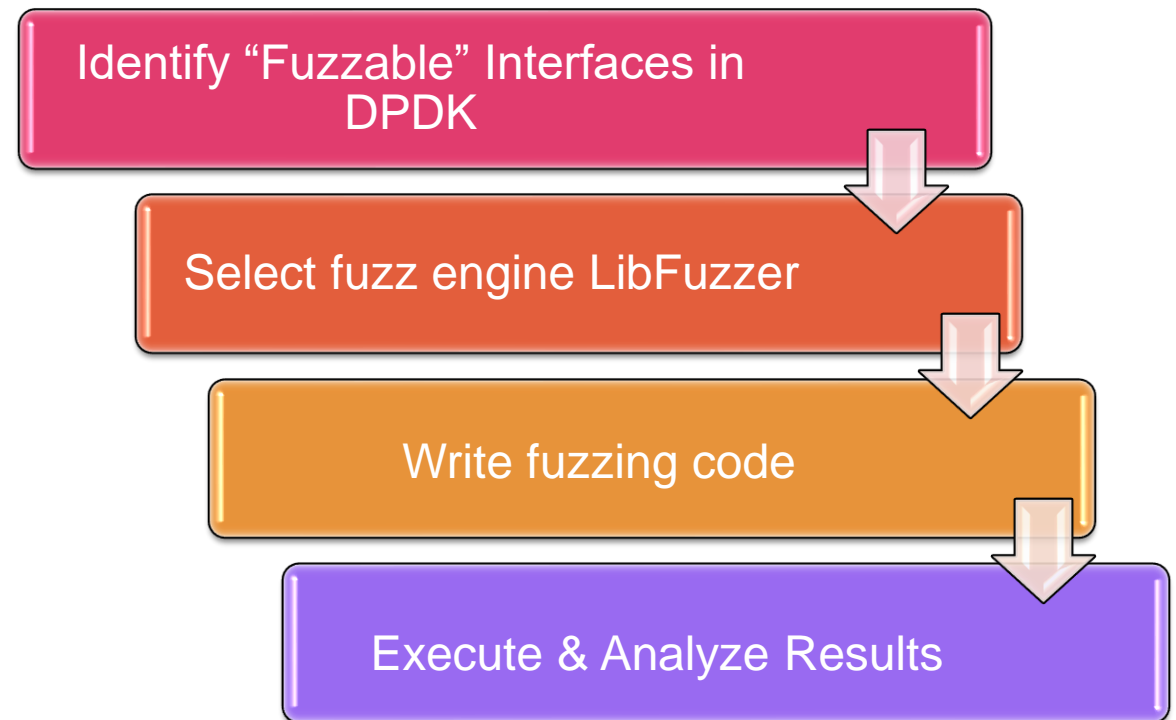
# How DPDK Fuzzing



## Fuzzing Definition

“Fuzzing is an automated **software testing technique** that involves providing invalid, unexpected, or **random** data as **inputs** to a computer program. The program is then **monitored** for exceptions such as **crashes**, **failing** built-in code **assertions**, or potential **memory leaks**”\*

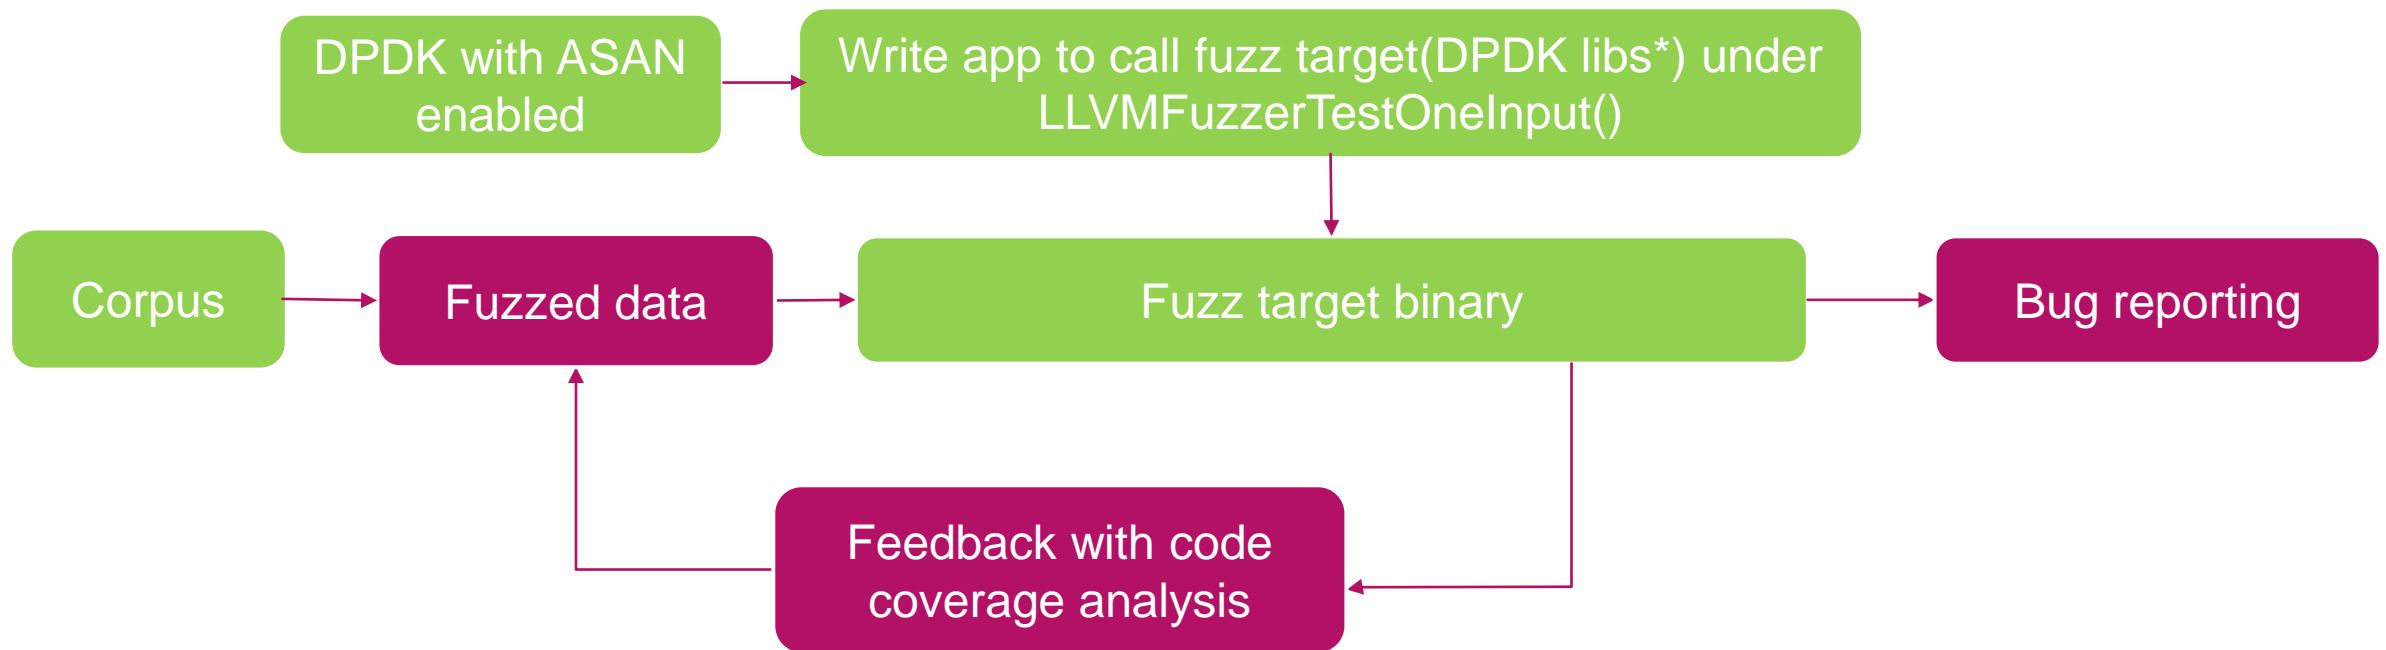
## DPDK Fuzzing Process



\*<https://en.wikipedia.org/wiki/Fuzzing>



# How DPDK Fuzzing



Note:  Need develop

 Libfuzzer framework

\* Note: DPDK Libs can include any number of libraries included in DPDK. Generally, would include e.g., distributor, fragmentation/reassembly, GRO/GSO,ETC.

# DPDK libFuzzer Deployment



- ❑ Build DPDK with ASAN enabled (Introduced in first part)
- ❑ Write app feeds fuzzed data into fuzz target under LLVMFuzzerTestOneInput()

```
//dpdk_fuzz.c  
  
int  
LLVMFuzzerTestOneInput(uint8_t *data, size_t size)  
{  
    do_IP_fragment(data, size);  
  
    return 0;  
}
```

- ❑ Build fuzz target

e.g., “cc=clang -g -fsanitize=address,fuzzer dpdk\_fuzz.c -o dpdk\_fuzz”

# DPDK libFuzzer Deployment



## ❑ DPDK initialize

```
root@dpdk-yinan-ntb1:/home/yinan/dpdk/dpdk-fuzzing-test-master/fuzzings/ip_fragmentation/build# ./ip_fragmentation /home/yinan/dpdk/corpus/
fuzzing> set arg -c/0x6/-n/4/-w/0000:03:00.0/--file-prefix=cy/---ip-type=ipv4
fuzzing> set arg -c/0x6/-n/4/-w/0000:03:00.0/--file-prefix=cy/--ip-type=ipv4
fuzzing> EAL: Detected 56 lcore(s)
EAL: Detected 2 NUMA nodes
EAL: WARNING! Base virtual address hint (0x100005000 != 0x7f61a6254000) not respected!
EAL: This may cause issues with mapping memory into secondary processes
EAL: Multi-process socket /var/run/dpdk/cy/mp_socket
EAL: Selected IOVA mode 'VA'
EAL: Probing VFIO support...
EAL: VFIO support initialized
```

## ❑ Begin fuzz

```
optind: 1 optopt:63 optarg:(null) argc:2 argv[0]:./ip_fragmentation argv[1]:--ip-type=ipv4
INFO: Seed: 2243594100
INFO: Loaded 1 modules (208 inline 8-bit counters): 208 [0x5b0e20, 0x5b0ef0),
INFO: Loaded 1 PC tables (208 PCs): 208 [0x572cf0, 0x5739f0),
INFO: 99 files found in /home/yinan/dpdk/corpus/
INFO: -max_len is not provided; libFuzzer will not generate inputs larger than 4096 bytes
INFO: seed corpus: files: 99 min: 1b max: 1365b total: 12438b rss: 44Mb
#100 INITED cov: 32 ft: 37 corp: 6/2174b exec/s: 0 rss: 46Mb
#101 NEW cov: 32 ft: 38 corp: 7/2341b lim: 1365 exec/s: 0 rss: 46Mb L: 167/1365 MS: 1 InsertRepeatedBytes-
#103 NEW cov: 34 ft: 40 corp: 8/2420b lim: 1365 exec/s: 0 rss: 46Mb L: 79/1365 MS: 2 EraseBytes-EraseBytes-
#111 NEW cov: 34 ft: 41 corp: 9/2612b lim: 1365 exec/s: 0 rss: 46Mb L: 192/1365 MS: 3 ShuffleBytes-InsertByte-CrossOver-
#120 NEW cov: 36 ft: 43 corp: 10/2751b lim: 1365 exec/s: 0 rss: 46Mb L: 139/1365 MS: 4 ChangeBit-CMP-CrossOver-ChangeBinInt- DE: "\x02\x00\x00\x00\x00\x00\x00\x00"-
#122 NEW cov: 37 ft: 44 corp: 11/2918b lim: 1365 exec/s: 0 rss: 46Mb L: 167/1365 MS: 2 ShuffleBytes-ChangeBinInt-
#130 NEW cov: 38 ft: 45 corp: 12/3968b lim: 1365 exec/s: 0 rss: 46Mb L: 1050/1365 MS: 3 CrossOver-EraseBytes-EraseBytes-
#141 NEW cov: 38 ft: 46 corp: 13/5253b lim: 1365 exec/s: 0 rss: 46Mb L: 1285/1365 MS: 1 CopyPart-
#294 REDUCE cov: 38 ft: 46 corp: 13/5238b lim: 1365 exec/s: 0 rss: 46Mb L: 1035/1365 MS: 3 EraseBytes-ChangeBit-InsertRepeatedBytes-
#404 REDUCE cov: 38 ft: 46 corp: 13/4268b lim: 1365 exec/s: 0 rss: 46Mb L: 315/1365 MS: 5 CopyPart-CrossOver-ChangeByte-ChangeASCIIInt-CrossOver-
#413 REDUCE cov: 38 ft: 46 corp: 13/3805b lim: 1365 exec/s: 0 rss: 46Mb L: 572/1365 MS: 4 InsertByte-InsertByte-PersAutoDict-EraseBytes- DE: "\x02\x00\x00\x00\x00\x00\x00\x00"-
```

# Thanks

[xueqin.lin@intel.com](mailto:xueqin.lin@intel.com) & [yinan.wang@intel.com](mailto:yinan.wang@intel.com)