



DPDK

—SUMMIT—

NORTH AMERICA

# Herd7 - Tool for memory model simulation

HONNAPPA NAGARAHALLI  
ARM

JULY 12-13, 2021

# Agenda

---

- Why do we need a tool?
- What to expect?
- Herd7
- Memory model simulation with Herd7
- Conclusion

# Why do we need a tool?

- Consider a simple thread synchronization problem – Message Passing
  - Writer thread – Update a data structure in shared memory
  - Reader thread – Read the data structure from shared memory
  - Writer uses a guard variable to communicate the readiness of data

```
/* Global variables */  
struct payload_s payload;  
int guard = 0;
```

```
/* Writer */  
/* Fill data */  
payload.a = 10;  
payload.b = 20;
```

A

```
/* Indicate payload is ready */  
atomic_store_explicit(&guard, 1, memory_order_relaxed);
```

B

2 events in Writer, A & B

```
/* Reader */  
/* Read the guard variable */  
if (atomic_load_explicit(&guard, memory_order_relaxed) != 0) {  
    /* payload is set by the writer, read payload */  
    have_fun(payload.a, payload.b);  
    .....  
}
```

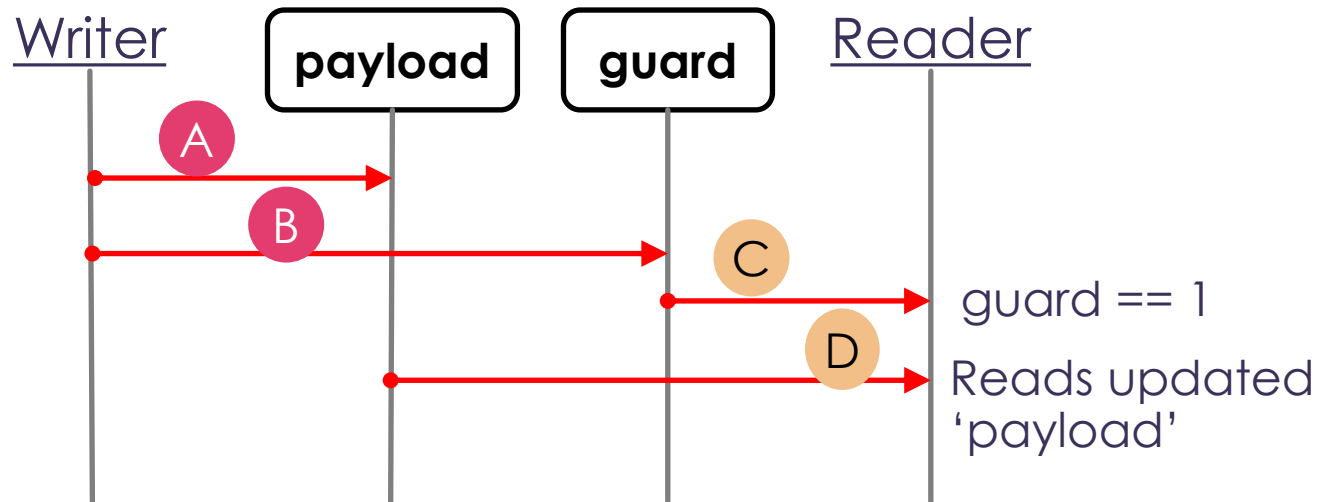
C

D

2 events in Reader, C & D

# Why do we need a tool?

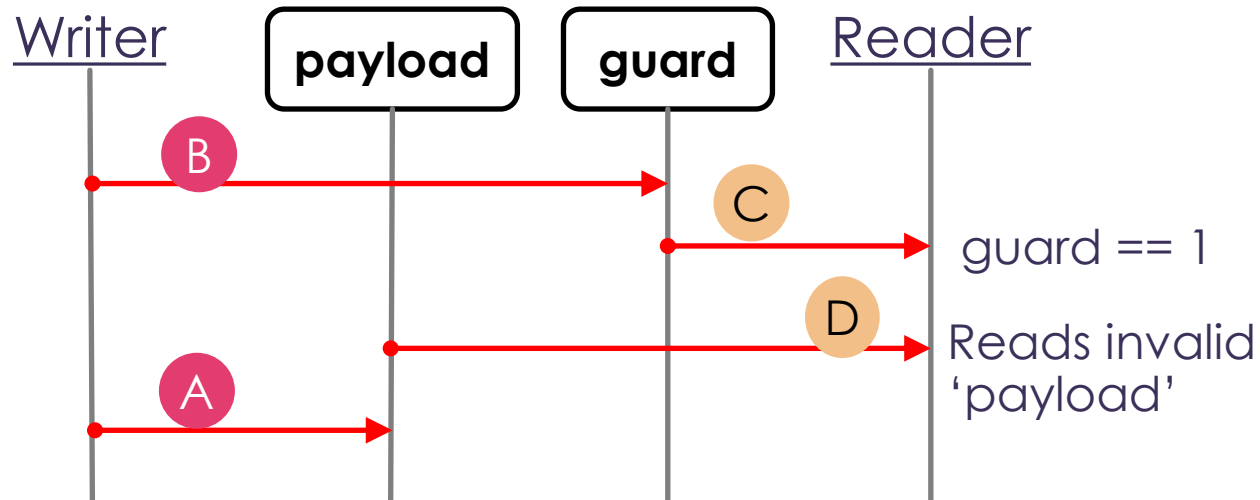
- In a machine with relaxed memory model, events can occur in any order
- Consider one such sequence ABCD



- No barriers required IF this sequence is executed

# Why do we need a tool?

- Consider another sequence BCDA



- This shows that 'A' and 'B' need to be execute in order, hence requiring barriers
- Similarly, if we consider sequence ADBC, 'C' and 'D' need to execute in order

# Why do we need a tool?

- So, what is the problem?
  - With 4 events, we have 24 possible sequences of events
  - Even with correct memory barriers – we have 6 valid sequences (ACBD, ACDB ...)
  - With 5 events – 120 sequences, 6 events – 720 sequences
  - It is not possible to write test cases to validate all these sequences
- Evolving CPU micro-architectures
  - CPU micro-architectures might implement a stronger memory model
  - Increasing performance requirements on CPUs results in more re-ordering
    - <https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/commit/?id=84a24bf8c52e66b7ac89ada5e3cfbe72d65c1896>
- Hence, we need a tool to validate the correctness of the algorithms

# What to expect

---

- The topic of Memory Model is complex
- I am a novice
- Goal
  - Introduce the tool
  - Reduce the barrier to get started
  - Simple examples, easy to understand, focus on the tool

- Part of herdtools7 Tool Chain
  - **herd7 – Memory model simulator**
  - diy7 / diycross7 / diyone7 – Litmus test generators
  - litmus7 / klitmus7 – Running the litmus tests on the actual hardware
- Tool page: <http://diy.inria.fr/>
- Documentation: <http://diy.inria.fr/doc/herd.html>
- Sources: <https://github.com/herd/herdtools7>
- Installing the tool on Linux
  - Cover in the backup slides
- Web Interface - <http://diy.inria.fr/www/>
  - C11 model will be enabled in the future



# Herd7 – Representing the algorithm – Litmus test

C Message-Passing

{ [payload] = 0; [guard] = 0; }

```
P0 (atomic_int* payload, atomic_int* guard) {
    atomic_store_explicit(payload, 2, memory_order_relaxed);
    atomic_store_explicit(guard, 1, memory_order_relaxed);
}

P1 (atomic_int* payload, atomic_int* guard) {
    rg = atomic_load_explicit(guard, memory_order_relaxed);
    if (rg == 1)
        rp = atomic_load_explicit(payload, memory_order_relaxed);
}
```

exists (1:rg=1 /\ 1:rp=0)

Name

Initialization

Programming Architecture

- Supports only small subset of C11 language
- Included – ‘if’, ‘while’, pointer dereferencing, simple expressions, atomic functions
- Excluded – ‘address-of’, compound types, function calls
- Supports Aarch64, PPC, RISC-V, x86

Program Code

- P0, P1 – Threads
- payload, guard – Global variables (shared memory)
- rg, rp – Local variables in thread1

Condition to assert

- Check if (rg ==1 && rp == 0) exists

# Herd7 – Analyzing the results

- `herd7 -c11 mp.litmus`

Test Message-Passing Allowed

States 4

1:rp=0; 1:rg=0;  
1:rp=0; 1:rg=1;  
1:rp=2; 1:rg=0;  
1:rp=2; 1:rg=1;

Shows the possible outcomes

Ok

Indicates condition exists

Witnesses

Positive: 1 Negative: 3

Condition exists (1:rg=1 /\ 1:rp=0)

Observation Message-Passing Sometimes 1 3

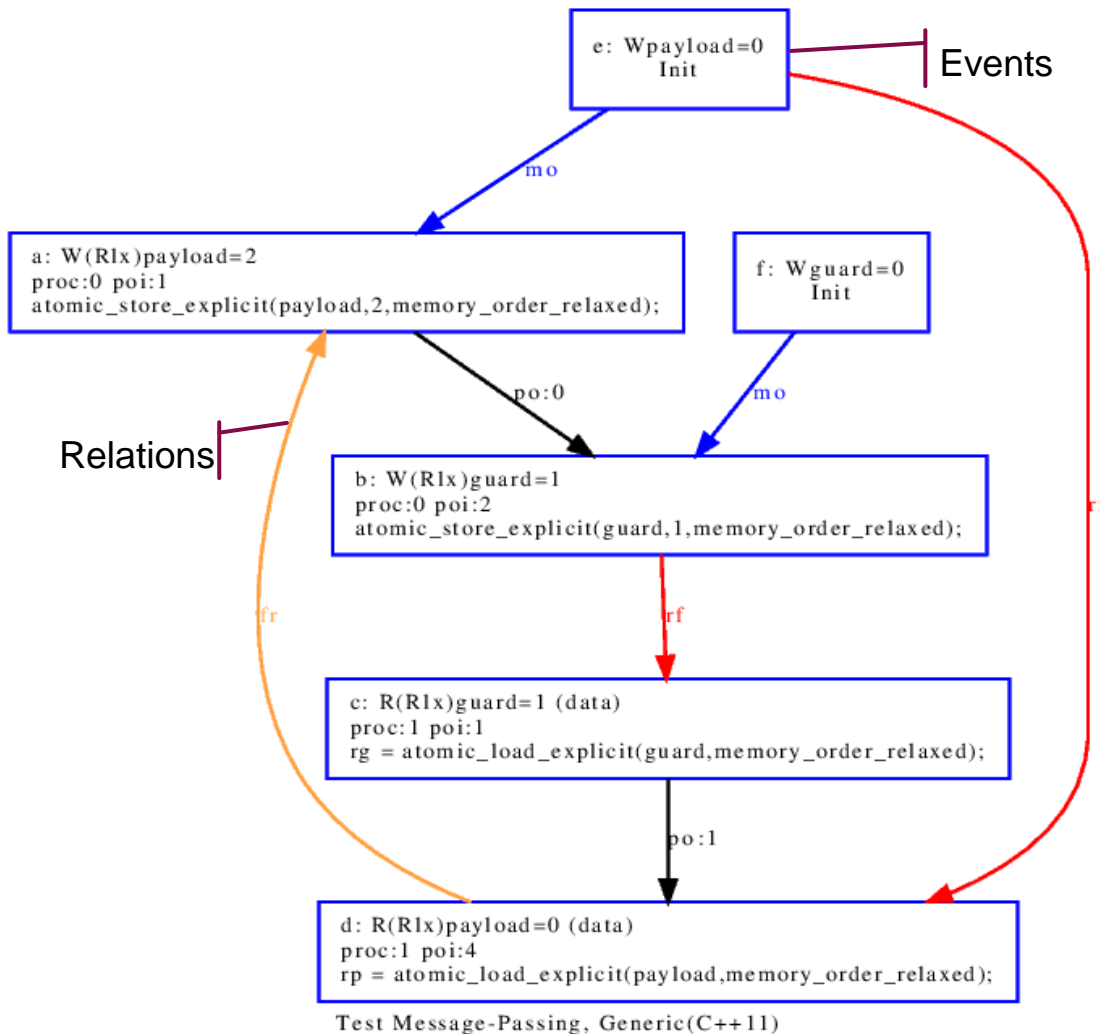
Indicates there is 1 sequence of events resulting in the condition.

Time Message-Passing 0.01

Hash=0d08978cd3e64ee878b665ee423d60df

# Herd7 – Fixing the algorithm

- herd7 -c11 -show prop -squished false -graph free -evince mp.litmus



- Events**
  - Wguard=0 : 0 is written to memory location 'guard'
  - Rpayload=0 : 0 is read from memory location 'payload'
  - proc:0 : Event in thread 0
- Relations**
  - po: Program Order: Order of the events in the code
  - rf: Read From: Which write event, the read event reads from
  - mo: Modification Order: Ordered history of all writes to a memory location
  - fr: From Read: Indicates all the writes that are overwriting a value read by the read event
- Identifying the problem**
  - Event 'R(Rlx)payload=0' reads from initialized value of 0
  - Event 'W(Rlx)payload=2' occurs after the read event
  - This sequence needs to be blocked

# Herd7 – Fixing the algorithm

C Message-Passing

```
{ [payload] = 0; [guard] = 0; }
```

```
P0 (atomic_int* payload, atomic_int* guard) {  
    atomic_store_explicit(payload, 2, memory_order_relaxed);  
    atomic_store_explicit(guard, 1, memory_order_release);  
}
```

```
P1 (atomic_int* payload, atomic_int* guard) {  
    rg = atomic_load_explicit(guard, memory_order_acquire);  
    if (rg == 1)  
        rp = atomic_load_explicit(payload, memory_order_relaxed);  
}
```

```
exists (1:rg=1 /\ 1:rp=0)
```

Test Message-Passing Allowed

States 2  
1:rg=0; 1:rp=0;  
1:rg=1; 1:rp=2;

No

Witnesses  
Positive: 0 Negative: 2  
Condition exists (1:rg=1 /\ 1:rp=0)  
Observation Message-Passing Never 0 2

Time Message-Passing 0.00  
Hash=a4450616bef403da2b0fc96f828429a2

- Other tools
  - CDS Checker – [http://plrg.eecs.uci.edu/software\\_page/42-2/](http://plrg.eecs.uci.edu/software_page/42-2/)
  - CPP Mem - <http://svr-pes20-cppmem.cl.cam.ac.uk/cppmem/index.html>
- Herdtools7
  - Supports architecture specific memory models
  - Does more exhaustive search for the re-ordering candidates (per available literature)
- Validate the DPDK algorithms using any of these tools

# Installing Herdtools7

---

- Environment used
  - Ubuntu 20.04.2
  - 5.4.0-77-generic
- Steps
  - Install Opam – **apt-get install opam**
  - Configure Opam – **opam init**
  - Install herdtools7 suite – **opam install herdtools7**
  - PDF viewer evince – Packaged in Ubuntu

# Questions?



DPDK

—SUMMIT—

NORTH AMERICA